

Реализация отказоустойчивого кластера PostgreSQL и проектирование базы данных



**Меня хорошо видно
& слышно?**



Защита проекта

Тема: Реализация отказоустойчивого кластера PostgreSQL и проектирование БД



Лонкин Дмитрий

DBA



Цель и задачи проекта

Цель проекта: спроектировать базу данных, построить отказоустойчивый кластер PostgreSQL.

1. Спроектировать структуру базы данных.
2. Построить отказоустойчивый кластер PostgreSQL
3. Реализовать систему резервного копирования
4. Система отчетов

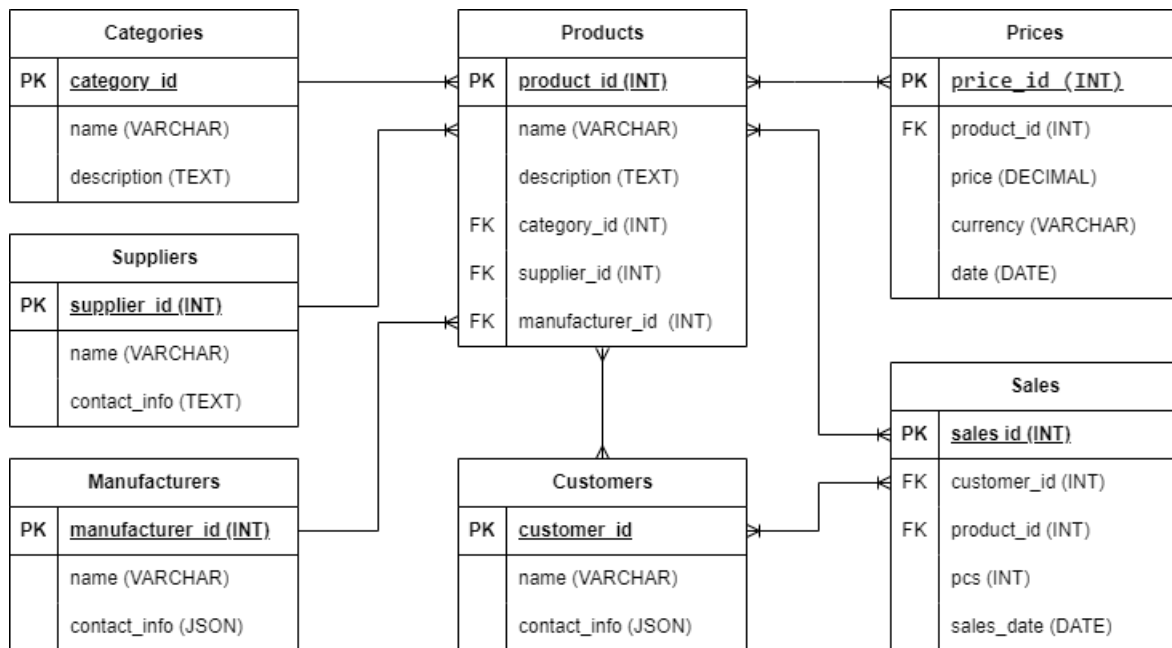
Какие технологии использовались

1. PostgreSQL под управлением patroni
2. Etcd
3. VipManager
4. Apache Superset
5. pgBackRest

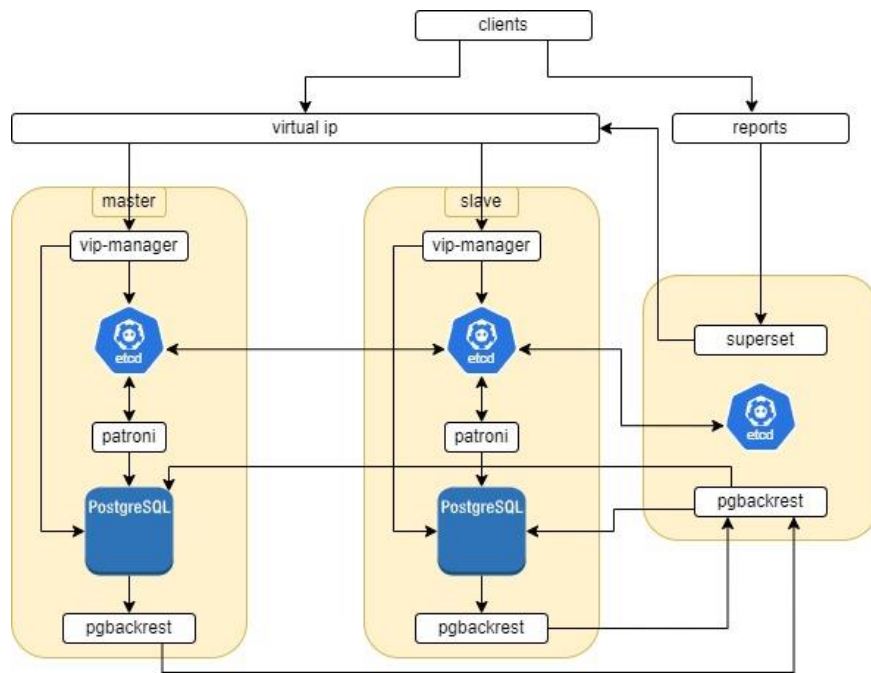


Что получилось

Структура БД



Архитектура системы



Кластер PostgreSQL / etcd

```
root@pg0-otus:~# patronictl -c /etc/patroni/config.yml list
```

```
+ Cluster: otus (7476070226399082672) -----+-----+-----+-----+
| Member | Host          | Role    | State    | TL | Lag in MB | Tags          |
+-----+-----+-----+-----+-----+-----+-----+
| pg0     | 192.168.88.65 | Leader  | running  | 7  |           | nosync: true |
| pg1     | 192.168.88.64 | Replica | streaming | 7  | 0         | nosync: true |
+-----+-----+-----+-----+-----+-----+-----+
```




```
pg1-otus postgres@internetshop=# select * from newtable ;
column1
-----
test2
(1 row)
```

```
pg1-otus postgres@internetshop=# select * from newtable ;
column1
-----
test2
(1 row)
```

```
pg1-otus postgres@internetshop=# insert into newtable values ('test1');
INSERT 0 1
```



```
pg1-otus postgres@internetshop=# select * from newtable ;
 column1
-----
 test2
 test1
(2 rows)
```

```
pg0-otus postgres@internetshop=# select * from newtable ;
 column1
-----
 test2
 test1
(2 rows)
```



Кластер etcd

```
root@mysql001-otus:~# etcdctl endpoint health --cluster
```

```
http://192.168.88.60:2379 is healthy: successfully committed proposal: took = 1.973148ms
```

```
http://192.168.88.65:2379 is healthy: successfully committed proposal: took = 2.599321ms
```

```
http://192.168.88.64:2379 is healthy: successfully committed proposal: took = 2.65843ms
```

```
root@mysql001-otus:~# etcdctl endpoint status --cluster --write-out=table
```

ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	IS LEARNER	RAFT TERM	RAFT INDEX	RAFT APPLIED INDEX	ERRORS
http://192.168.88.64:2379	1c938984d2e351c2	3.4.30	164 kB	false	false	91887	557	557	
http://192.168.88.60:2379	6ea35c095e2a522b	3.4.30	164 kB	false	false	91887	557	557	
http://192.168.88.65:2379	7543f4a133d1e671	3.4.30	164 kB	true	false	91887	557	557	



etcd

```
root@pg1-otus:~# etcdctl get / *
/otus/otus/config
{"archive_command":"pgbackrest --stanza=ishop archive-push
%p","log_directory":"/pg_audit","logging_collector":true,"log_truncate_on_rotation":true,"log_rotation_age":"1d","log_rotation_size":0,"
log_line_prefix":"%t [%p-%l] %r %q%u@d ","log_filename":"postgresql-%Y-%m-%d-
%H%M.log","loop_wait":10,"maximum_lag_on_failover":1048576,"retry_timeout":10,"ttl":30}
/otus/otus/failover
{}
/otus/otus/history
[[1,67109024,"no recovery target specified","2025-02-27T15:19:38.582177+03:00","pg0"],[2,117440672,"no recovery target specified","2025-
02-27T17:08:31.591715+03:00","pg1"],[3,234881184,"no recovery target specified","2025-02-
27T17:51:27.395146+03:00","pg0"],[4,285212832,"no recovery target specified","2025-02-27T18:42:31.398089+03:00","pg1"],[5,352336560,"no
recovery target specified","2025-02-28T14:38:56.147314+03:00","pg0"],[6,352500336,"no recovery target specified","2025-03-
04T21:13:38.616000+03:00","pg0"]]
/otus/otus/initialize
7476070226399082672
/otus/otus/leader
pg0
/otus/otus/members/pg0
{"conn_url":"postgres://192.168.88.65:5432/postgres","api_url":"http://192.168.88.65:8008/patroni","state":"running","role":"primary","v
ersion":"4.0.4","tags":{"nosync":true},"xlog_location":469762800,"timeline":7}
/otus/otus/members/pg1
{"conn_url":"postgres://192.168.88.64:5432/postgres","api_url":"http://192.168.88.64:8008/patroni","state":"running","role":"replica","v
ersion":"4.0.4","tags":{"nosync":true},"xlog_location":469762800,"replication_state":"streaming","timeline":7}
/otus/otus/status
{"optime":469762800,"slots":{"pg1":469762800,"pg0":469762800},"retain_slots":["pg0","pg1"]}
```



VipManager

```
root@pg0-otus:~# cat /etc/default/vip-manager.yml
interval: 1000
trigger-key: "/otus/otus/leader"
trigger-value: "pg0"
ip: 192.168.88.123 # the virtual ip address to manage
netmask: 24 # netmask for the virtual ip
interface: enp0s3 #interface to which the virtual ip will be added
hosting-type: basic # possible values: basic, or hetzner.
dcs-type: etcd # etcd, consul or patroni
dcs-endpoints:
  - http://192.168.88.64:2379
  - http://192.168.88.65:2379
retry-num: 2
retry-after: 250 #in milliseconds
verbose: false
```



```
root@pg1-otus:~# cat /etc/default/vip-manager.yml
interval: 1000
trigger-key: "/otus/otus/leader"
trigger-value: "pg1"
ip: 192.168.88.123 # the virtual ip address to manage
netmask: 24 # netmask for the virtual ip
interface: enp0s3 #interface to which the virtual ip will be added
hosting-type: basic # possible values: basic, or hetzner.
dcs-type: etcd # etcd, consul or patroni
dcs-endpoints:
  - http://192.168.88.64:2379
  - http://192.168.88.65:2379
retry-num: 2
retry-after: 250 #in milliseconds
verbose: false
```



Система резервного копирования

```
root@pgbackrest-otus:~# cat /etc/pgbackrest.conf
```

```
[global]
```

```
process-max=6
```

```
repol-path=/pgbackrest
```

```
repol-retention-full-type=time
```

```
repol-retention-full=3
```

```
repol-retention-diff=7
```

```
start-fast=y
```

```
log-level-file=detail
```

```
config-include-path=/etc/pgbackrest/conf.d
```

```
[global:archive-push]
```

```
compress-level=3
```

```
[ishop]
```

```
pg1-host=192.168.88.65
```

```
pg1-path=/pg_data
```

```
pg1-socket-path=/var/run/postgresql
```

```
pg2-host=192.168.88.64
```

```
pg2-path=/pg_data
```

```
pg2-socket-path=/var/run/postgresql
```




Расписание резервного копирования

```
root@pgbackrest-otus:~# crontab -l  
# m h dom mon dow    command  
0 23 * * 0 sudo -u pgbackrest pgbackrest --stanza=ishop --type=full backup  
0 23 * * 1-6 sudo -u pgbackrest pgbackrest --stanza=ishop --type=diff backup
```



Superset

Time Series Line Chart 2024



Total quantity for day

purchase_date		total_quantity	total_amount
2024-10-23		8	\$3368.08
2024-10-22		20	\$15447.38
2024-10-19		8	\$1998.68
2024-10-18		7	\$3825.68
2024-10-16		1	\$657.25
2024-10-15		9	\$6505.38
2024-10-14		20	\$12205.98
2024-10-13		11	\$2890.25
2024-10-10		5	\$4329.90
2024-10-09		18	\$8318.05



Выводы

1. Цель проекта достигнута (кластер, БД, РК, Отчеты)
2. Есть возможность добавить PostgreSQL для отчета.

Вопросы и рекомендации



если есть вопросы



если вопросов нет

Спасибо за внимание!

